

genAI

[generative AI]

Good old data mining ['GODaM' :)]...

So far, we have studied many data mining algorithms (including NNs), which all learn (model) patterns (between outputs and inputs) in existing data, and use that to classify/calculate new outputs based on new inputs.

Given existing data: existing inputs → existing outputs

Learn (model) the pattern in the data

Use the model to compute: NEW input → ? NEW output

So... what's this 'genAI' thing?

'Generative AI' - a revolution

Generative AI, very loosely speaking, 'runs a neural network BACKWARDS'!

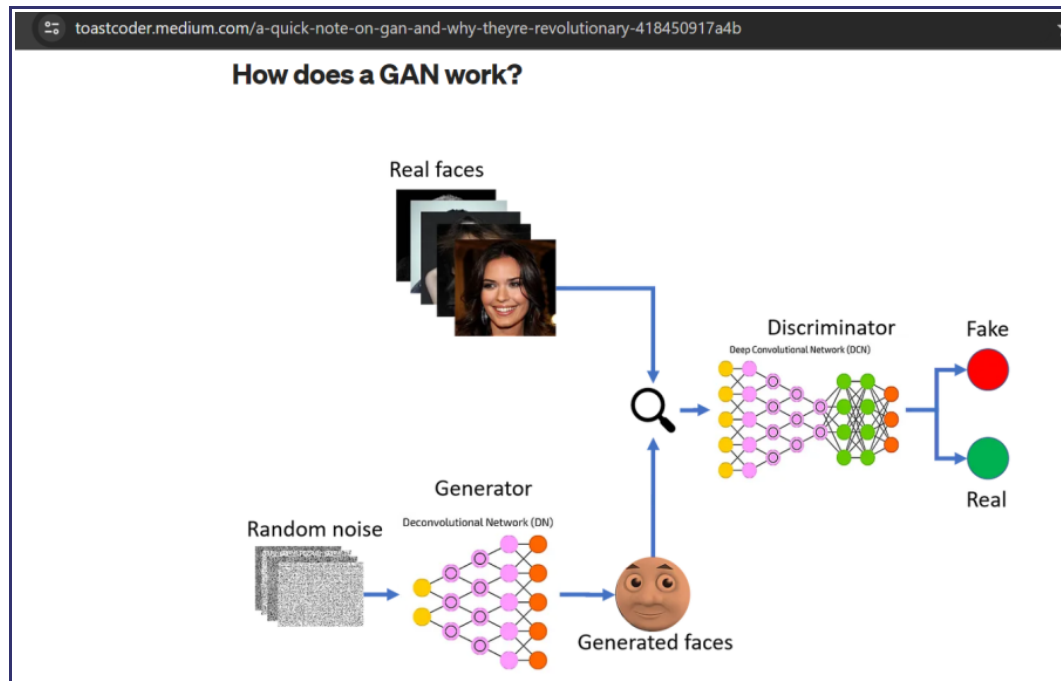
Rather than learn to classify new data using existing data, why not GENERATE new data instead?

Researchers tried this, but with unimpressive results.

In 2014, Ian Goodfellow got a much better idea than the 'SOTA' - why not pair up TWO NNs in opposing order - one a generator (eager 'student'), and the other, a discriminator (strict 'teacher')? His invention is called a 'GAN'.

GAN

GAN stands for Generative Adversarial Network



More: https://developers.google.com/machine-learning/gan/gan_structure and <https://machinelearningmastery.com/what-are-generative-adversarial-networks-gans/>

Style transfer

An early use of genAI was/is to "add style" to imagery: <https://arxiv.org/abs/1508.06576>

anmol19005.medium.com/evolution-of-style-transfer-techniques-1a892e796475


Content ImageStyle ImageOutput Image

The Deep Learning magic

Our starting point was the original algorithm for neural style transfer by [Gatys et al \(2015\)](#) based on optimizing an image to match the content and style of another image by manipulating some clever losses.

In summary, the general idea is:

- Take an image for style (`style_img`)
- Take an image for content (`content_img`)
- Take a pre-trained neural net
- Create an “output” image where each pixel is a parameter to optimize
- Optimize the “output” image so that the content loss and style loss are minimized

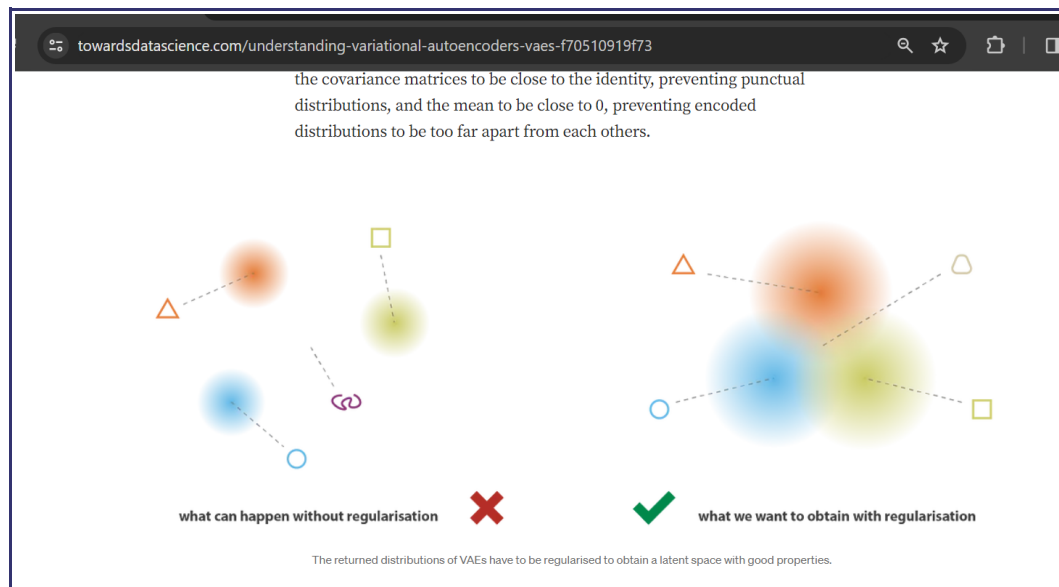


Encoders/decoders, autoencoders, VAEs

Encoder: a function (NN) that maps original input to LATENT/ENCODED space [decoder is the reverse]

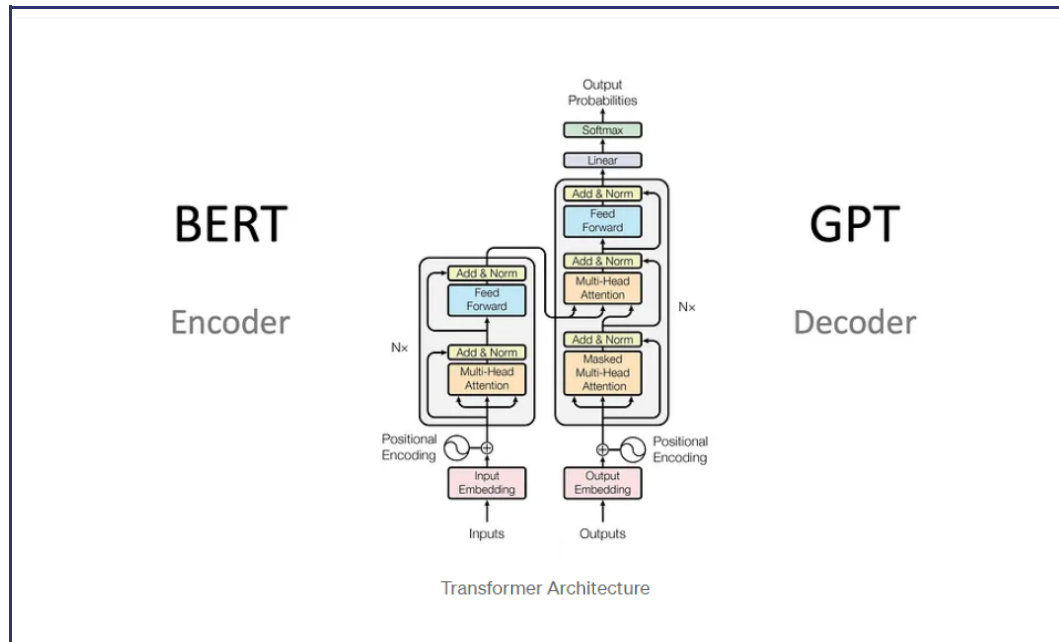
Autoencoder: encoder + decoder combination - can GENERATE NEW OUTPUT (by interpolating a random point in latent space)!

Variational AE: the encoder produces a distribution rather than a single point [and the decoder uses a sampled point from the distribution].



Transformers

'Attention is all you need': <https://arxiv.org/abs/1706.03762>



Revolution: non-fixed size and parallelizable self-attention mechanism (ie. computing word affinities).

More:

- <https://jalammar.github.io/illustrated-transformer/>
- <https://medium.com/@yulemoon/detailed-explanations-of-transformer-step-by-step-dc32d90b3a98>
- https://www.youtube.com/watch?v=zjkBMFhNj_g

The decoder takes a prompt (new point in latent space), INTERPOLATES over inputs (ALL English!!), generates output.

Transformers to... ChatGPT!

Here is a great explanation of ChatGPT: <https://writings.stephenwolfram.com/2023/02/what-is-chatgpt-doing-and-why-does-it-work/>

The core pre-trained LLM needs to be augmented with HFRL [a form of fine-tuning], to produce acceptable responses.

GPT extensions

- multimodal (eg text, images)
- plugin API
- GPT Store
- LLM apps, eg. <https://blog.llamaindex.ai/introducing-llama-packs-e14f453b913a> [and <https://llamahub.ai/>]

LLM extensions

- larger context [eg. RMT: <https://arxiv.org/abs/2304.11062>, and <https://hazyresearch.stanford.edu/blog/2023-03-07-hyena> and <https://bdtechtalks.com/2023/11/27/streamingllm/amp/>]
- infinite memory! [https://medium.com/@jordan_gibbs/how-to-create-your-own-gpt-voice-assistant-with-infinite-chat-memory-in-python-d8b8e93f6b21]
- architecture alterations (eg Rethinking Attention: <https://arxiv.org/abs/2311.10642>), alternate position encodings [including NO: <https://arxiv.org/pdf/2203.16634.pdf>, context-aware, rotary encoding...]
- open source!
- quantization [of weights to 16/8...bits, to compress model size]
- SLMs! [latest: <https://mistral.ai/>, <https://starling.cs.berkeley.edu/>, Orca 2...] (almost all at <https://huggingface.co/models>)

Prompt extensions/ LLM frameworks

- LangChain
- LlamaIndex
- Haystack
- <https://cassio.org/>
- ...
- CoT [chain of thoughts], ToT [tree of thoughts]...
- LangChain Expression Language [LCEL]
- agents!

Response extensions

- fine-tuning, eg. <https://magazine.sebastianraschka.com/p/practical-tips-for-finetuning-llms> and <https://www.databricks.com/blog/efficient-fine-tuning-lora-guide-llms>
- **RAG!!** Two kinds (eg. <https://ai.plainenglish.io/beyond-tables-and-vectors-knowledge-graphs-for-ai-reasoning-46f0f8721894>), more kinds... [eg. <https://artificialcorner.com/ive-created-a-custom-gpt-that-scrapes-data-from-websites-9086aff58105>]
- vector DBs https://medium.com/@zilliz_learn/what-is-a-real-vector-database-b391b0468279 and <https://tiledb.com/blog/why-tiledb-as-a-vector-database>
- LMSQL! <https://towardsdatascience.com/lmql-sql-for-language-models-d7486d88c541>

Custom GPTs

As a result of the extensions listed above, there is bound to be numerous, narrow-purpose GPTs, eg. <https://chat.openai.com/g/g-kCfSC3b10-analystgpt> - it's in a way, back to 'expert systems' AI from the mid-80s and early 90s :)

Here is an architecture useful for building at-scale RAG apps: <https://www.pinecone.io/learn/aws-reference-architecture/>

genAI: other (non-plaintext) content

- CODE!!!
- images
- video
- music
- 3D CG
- ...

Issues?

<https://www.theguardian.com/books/2023/nov/15/hallucinate-cambridge-dictionary-word-of-the-year>

GenAI Against Humanity: <https://arxiv.org/pdf/2310.00737.pdf>